



# Machine Learning Capstone Project

**John Gimlin**

jdgimlin@gmail.com

09/01/2020



# Project Overview

For this capstone project, codecademy provided a dataset from **okcupid**, an online dating app that matches users based on a number of multiple choice and short-answer essay questions.

*From codecademy:*

***“The purpose of this capstone is to practice formulating questions and implementing Machine Learning techniques to answer those questions.”***



# Project Tasks



1. Explore the Data
2. Visualize the Data
3. Formulate a Question
4. Augment the Data
5. Compare Classification Models
6. Compare Regression Models
7. Analyze Model Metrics
8. Conclusion





# 1. Explore the Data



# 1. Explore the Data

The okcupid dataset consists of:

- **59,946** rows and **31** columns
- **3** numerical columns
- **28** object data-type columns

## Columns

['age', 'body\_type', 'diet', 'drinks', 'drugs', 'education', 'essay0', 'essay1', 'essay2', 'essay3', 'essay4', 'essay5', 'essay6', 'essay7', 'essay8', 'essay9', 'ethnicity', 'height', 'income', 'job', 'last\_online', 'location', 'offspring', 'orientation', 'pets', 'religion', 'sex', 'sign', 'smokes', 'speaks', 'status']

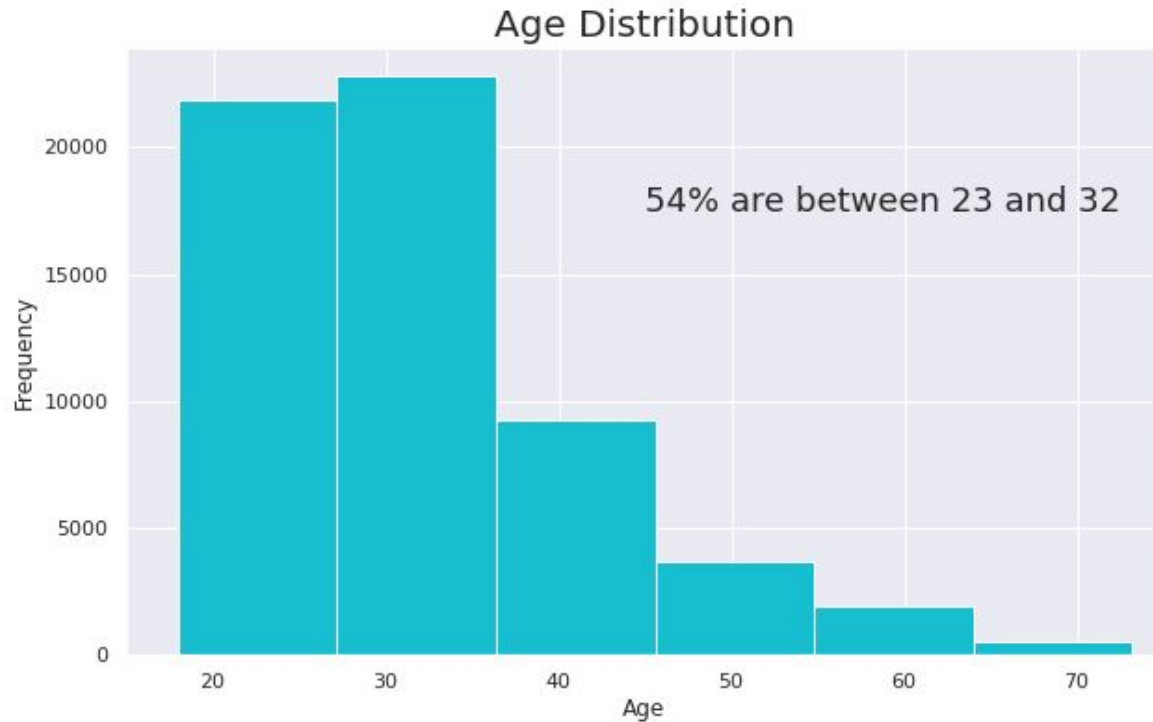




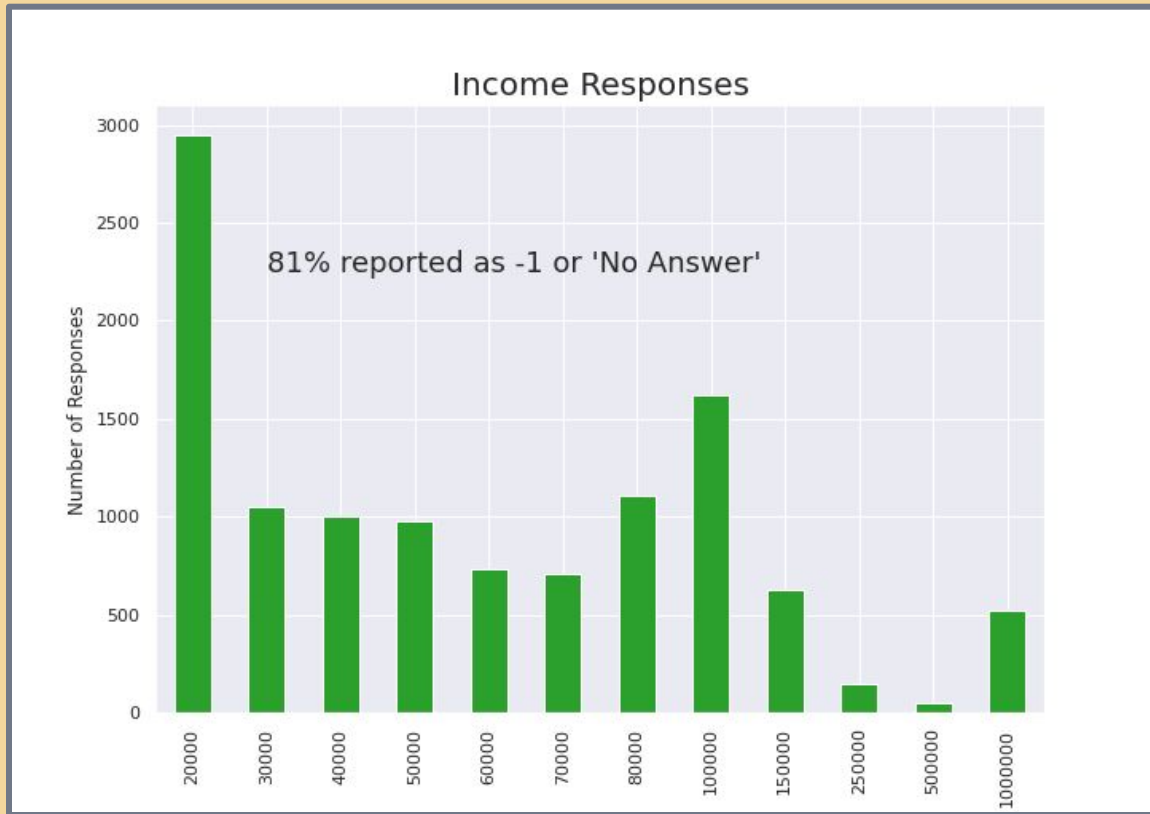
## 2. Visualize the Data



# Visualize the Data



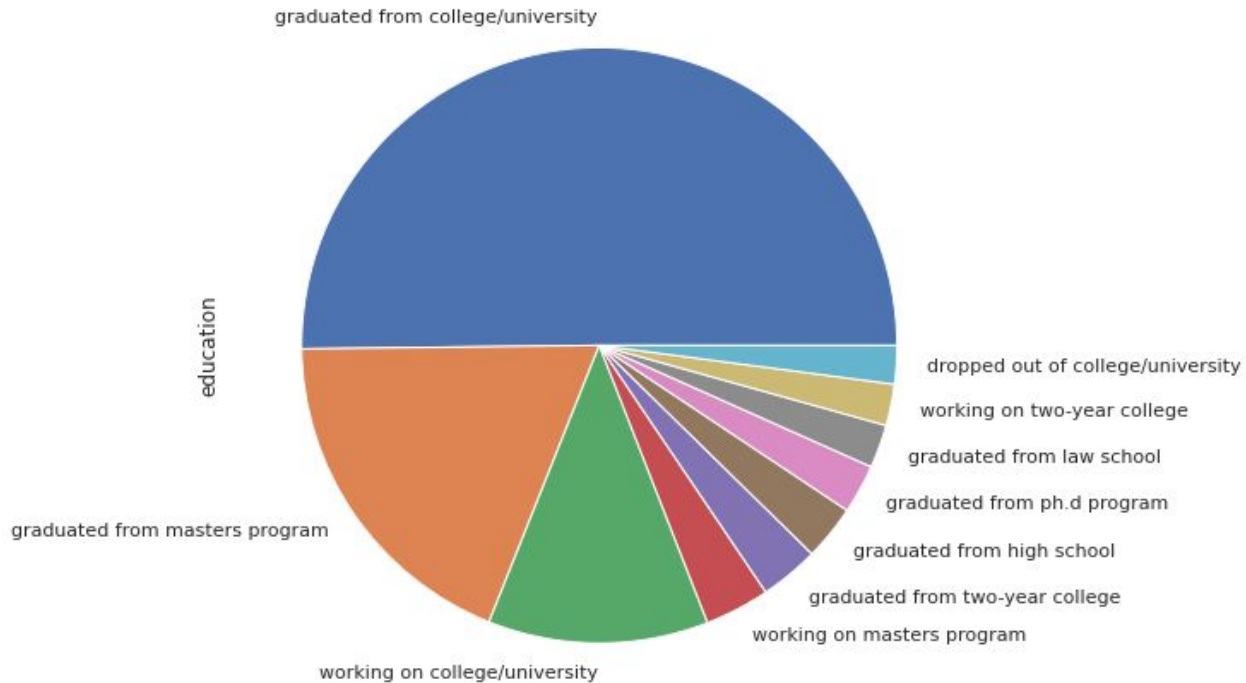
# Visualize the Data






# Visualize the Data

Top 10 Education Responses



# Visualize the Data





# 3. Formulate a Question



### 3. Formulate a Question

**Classification Model Question:** Can we predict **gender** based on the following features:

1. Whether a user prefers cats or dogs from the **pets** column
2. A user's **body\_type** and **height**
3. How much a user **drinks** or **smokes**

I chose to predict gender because it's a binary, categorical response. After trying many different features, I chose the above out of general interest and possible correlation to the response category.



### 3. Formulate a Question

**Regression Model Question:** Can we predict **height** based on the following features:

1. A user's **diet**
2. Their **drug** use
3. Their gender, or **sex**

I was limited to only three possible regression response columns. I didn't choose **income** due to the high number of null values. That left **age** or **height**. I chose **height** because it had more correlation to the feature columns.






## 4. Augment the Data



## 4. Augment the Data

### **A note on classifying categorical data with numeric values:**

- Codecademy does not cover any preprocessing techniques or tools, such as dummy variables with pandas or scikit learn's OneHotEncoder. As a result I mapped categorical responses into new columns using integers based solely on common sense and instinct.
- This was complicated by the fact the okcupid survey contained many possible answers for essentially the same response across all of the categorical response questions.
- As a project requirement the next slide will show how I chose to map some of the feature categories used in the modeling. This method was used for the classification and regression models.



## 4. Augment the Data

Sample mapping of data into new columns:

```
# Map "pets"
# Create & map feature column: pets_code

pets_mapping = {
    "likes dogs and likes cats": 1,
    "likes dogs": 0,
    "likes dogs and has cats": 2,
    "has dogs": 0,
    "has dogs and likes cats": 0,
    "likes dogs and dislikes cats": 0,
    "has dogs and has cats": 1,
    "has cats": 2,
    "likes cats": 2,
    "has dogs and dislikes cats": 0,
    "dislikes dogs and likes cats": 2,
    "dislikes dogs and dislikes cats": -1,
    "dislikes cats": -1,
    "dislikes dogs and has cats": 2,
    "dislikes dogs": -1}
df_class["pets_code"] = df_class.pets.map(pets_mapping)
```

```
# Map "body_type"
# Create & map feature column: body_type_code

bt_mapping = {
    "average": 0,
    "fit": 1,
    "athletic": 1,
    "thin": -1,
    "curvy": 2,
    "a little extra": 2,
    "skinny": -1,
    "full figured": 2,
    "overweight": 3,
    "jacked": 1,
    "used up": -1,
    "rather not say": 0}
df_class["body_type_code"] = df_class.body_type.map(bt_mapping)
```



## 4. Augment the Data

Sample mapping of data into new columns:

```
# Map "drinks"
# Create & map feature column: drinks_code

drinks_mapping = {
    "socially": 0,
    "rarely": 1,
    "often": 2,
    "not at all": -1,
    "very often": 3,
    "desperately": 3}
df_class["drinks_code"] = df_class.drinks.map(drinks_mapping)
```

```
# Map "smokes"
# Create & map feature column: smokes_code

smokes_mapping = {
    "no": 0,
    "sometimes": 1,
    "when drinking": 1,
    "yes": 2,
    "trying to quit": 2,}
df_class["smokes_code"] = df_class.smokes.map(smokes_mapping)
```

```
# Map response column y: "sex"
# Create new column: sex_code

df_class["sex_code"] = df_class.loc[:, "sex"].apply(lambda x: 1 if x == "m" else 0)
```

## 4. Augment the Data (Classification Model)

New column created: `sex_code`

This served as the response column where “male” was mapped as 1 and “female” as 0.

```
df_class.sex_code.value_counts()
```

```
1    19900  
0    14522
```

```
Name: sex_code, dtype: int64
```

New column created: `pets_code`

Mapped user preferences for cats, dogs, both, and neither as numerical values.

```
df_class.pets_code.value_counts()
```

```
1    14122  
0    13945  
2     6026  
-1     329
```

```
Name: pets_code, dtype: int64
```

New column created: `body_type_code`

Mapped user responses for thin, average, fit, etc. as numerical values.

```
df_class.body_type_code.value_counts()
```

```
1    14888  
0     9465  
2     5520  
-1    4209  
3      340
```

```
Name: body_type_code, dtype: int64
```

## 4. Augment the Data (Classification Model) cont.

New column created: `drinks_code`

Mapped user responses for drinking alcohol into numerical values.

New column created: `smokes_code`

Mapped user responses for smoking as numerical values.

```
df_class.drinks_code.value_counts()  
0      24907  
1       3801  
2       3091  
-1      2136  
3        487  
Name: drinks_code, dtype: int64
```

```
df_class.smokes_code.value_counts()  
0      27538  
1       4427  
2       2457  
Name: smokes_code, dtype: int64
```

## 4. Augment the Data

### Classification Model

Unneeded columns were dropped and rows with NaN values were removed. This left a dataframe with 34,422 rows and 11 columns.

```
df_class.sample(5, random_state=40)
```

	pets	body_type	height	drinks	smokes	sex	sex_code	pets_code	body_type_code	drinks_code	smokes_code
38235	likes dogs and likes cats	curvy	60.0	socially	no	f	0	1	2	0	0
56177	likes dogs	thin	62.0	socially	no	f	0	0	-1	0	0
55224	likes dogs	fit	68.0	socially	no	m	1	0	1	0	0
27481	likes dogs and likes cats	skinny	70.0	often	when drinking	m	1	1	-1	2	1
59485	has cats	average	77.0	often	yes	f	0	2	0	2	2



## 4. Augment the Data (Regression Model)

New column created: `diet_code`

Mapped user diet preference, ie “vegetarian”, “vegan”, etc. as numerical values.

New column created: `drugs_code`

Mapped user responses for drug use as numerical values.

New column created: `sex_code`

Mapped male users as 1 and female users as 0.

```
df_reg.diet_code.value_counts()
```

```
1    20617
2     3988
0     1313
3       149
Name: diet_code, dtype: int64
```

```
df_reg.drugs_code.value_counts()
```

```
0    21041
1     4758
2       268
Name: drugs_code, dtype: int64
```

```
df_reg.sex_code.value_counts()
```

```
1    15512
0    10555
Name: sex_code, dtype: int64
```



## 4. Augment the Data

### Regression Model

Unneeded columns were dropped and rows with NaN values were removed. This left a dataframe with 26,067 rows and 11 columns.

```
df_reg.sample(5, random_state=4)
```

	diet	drinks	drugs	smokes	sex	height	diet_code	drugs_code	sex_code	drinks_code	smokes_code
28848	mostly anything	socially	never	no	m	72.0	1	0	1	1	0
45549	anything	socially	never	no	m	75.0	1	0	1	1	0
53378	kosher	socially	never	no	f	65.0	3	0	0	1	0
31109	anything	socially	never	no	m	68.0	1	0	1	1	0
53849	mostly vegetarian	socially	never	no	f	67.0	2	0	0	1	0



# 5. Compare Classification Models



## 5. Compare Classification Models


I used **Logistic Regression** and **K Nearest Neighbors (K=5)** for the classification models.

Both returned a scoring accuracy of approximately **83%** on the training set .

Logistic Regression runs notably faster: **62 ms** versus **198 ms** for KNN [K=5].







# 6. Compare Regression Models

## 6. Compare Regression Models

I used **Multiple Linear Regression** and **K Neighbors Regressor [K=5]** for the regression models.

MLR returned a coefficient of determination [ $R^2$ ] of 44% on the training set. KNR [K=5] scored 40%

MLR runs notably faster: 11 ms versus 253 ms for KNR [K=5].





# 7. Analyze Model Metrics



## 7. Analyze Model Metrics: Classification Models

For model scoring I used scikit learn's Train Test Split utility to create a testing set consisting of 25% of the available data.

### K Nearest Neighbors K=5

Accuracy: 0.818  
Precision: 0.817  
Recall: 0.818

### Logistic Regression

Accuracy: 0.827  
Precision: 0.826  
Recall: 0.827


The **null accuracy** [score from choosing the most common class] is **58%** so the models each did pretty well, comparatively.

### Null Accuracy

```
# Null Accuracy  
y.value_counts(normalize=True)  
  
1    0.578119  
0    0.421881  
Name: sex_code, dtype: float64
```

## 7. Analyze Model Metrics: Classification Models

I checked to see if I could find a better value for K in the KNN model.



K=18 obtained a slightly higher model accuracy score of 0.834 around 1.6 percentage points higher than K=5.

## 7. Analyze Model Metrics: Regression Models

For model scoring I used scikit learn's Train Test Split utility to create a testing set consisting of 25% of the available data.

### Multiple Linear Regression

Training set  $R^2$ : 0.440  
Testing set  $R^2$ : 0.449  
RMSE: 2.916

### K Neighbors Regressor

Training set  $R^2$ : 0.404  
Testing set  $R^2$ : 0.413  
RMSE: 2.929

The coefficient of determination [ $R^2$ ] for both models show some correlation between the feature class and the response class, **height**.

RMSE of approx. **2.9** indicates error is within three units of the y axis, or in this case, inches.

### Note:


The project requirements called for accuracy and precision scores for regression models, however:

“Evaluation metrics for classification problems, such as **accuracy**, are not useful for regression problems. Instead, we need evaluation metrics designed for comparing continuous values, like the Root Mean Squared Error, or **RMSE**.”

Source: [DataSchool](https://datacamp.com/learn/machine-learning-with-python/part-4/linear-regression-with-python/)

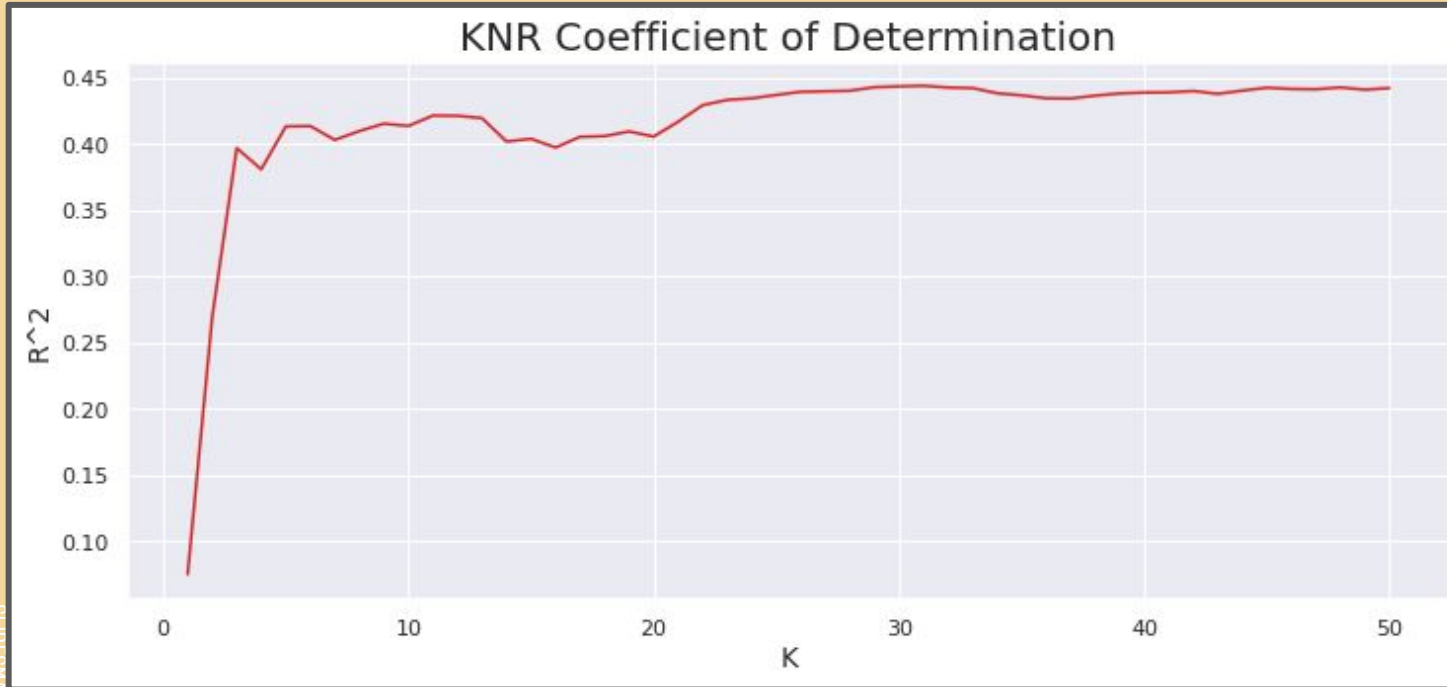
## 7. Analyze Model Metrics: Regression Models

### Multiple Linear Regression



## 7. Analyze Model Metrics: Regression Models

I checked to see if I could find a better value for K in the KNR model.



K=31 obtained a slightly higher model  $R^2$  score of 0.444 only 3.1 percentage points higher than K=5.





# 8. Conclusion



## 8. Conclusion

Bottom line is that I was able to construct two classification models that were decently predictive and two regression models that showed a faint linear relationship between the feature and response classes.

The provided dataset was honestly terrible, in my opinion. This may have been intentional from a teaching standpoint, but it wasn't very conducive for machine learning.

As it was data from  $\approx 60k$  users on a dating site, it's possible none of the data was truthful. This was apparent looking at the **income** values [surprising number of millionaires!] and reported **heights** [amazing number of people under 4-feet and over 7-feet tall]. Shocking it included responses from two users over the **age** of 108! And finally, the **10 essay** columns were utterly useless and made up  $\approx 30\%$  of the dataset.

In the end I did learn a lot about cleaning and mapping data for machine learning. I just wish the data provided would have had greater validity and been more interesting to construct machine learning models with.

